

```
In [2]: import pandas as pd
import numpy as np
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import classification_report
```

```
In [4]: # Load the CSV file (since it's in the same folder)
df = pd.read_csv('StudentsPerformance.csv')

# Preview the dataset
df.head()
```

```
Out[4]:
```

	gender	race/ethnicity	parental level of education	lunch	test preparation course	math score	reading score	writing score
0	female	group B	bachelor's degree	standard	none	72	72	74
1	female	group C	some college	standard	completed	69	90	88
2	female	group B	master's degree	standard	none	90	95	93
3	male	group A	associate's degree	free/reduced	none	47	57	44
4	male	group C	some college	standard	none	76	78	75

```
In [6]: # Label encode categorical columns
label_encoders = {}
for col in df.select_dtypes(include='object').columns:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    label_encoders[col] = le
```

```
In [8]: # Calculate average score
df['average_score'] = df[['math score', 'reading score', 'writing score']].mean(axis=1)

# Create binary classification target
df['pass'] = (df['average_score'] >= 60).astype(int)

# Drop the average column to avoid data Leakage
df.drop(columns=['average_score'], inplace=True)
```

```
In [10]: # Separate features and target
X = df.drop('pass', axis=1)
y = df['pass']
```

```
# Split into train and test sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Standardize features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
In [14]: lr_model = LogisticRegression(max_iter=1000, random_state=42)
lr_model.fit(X_train_scaled, y_train)
y_pred_lr = lr_model.predict(X_test_scaled)
```

```
In [16]: knn_model = KNeighborsClassifier(n_neighbors=5)
knn_model.fit(X_train_scaled, y_train)
y_pred_knn = knn_model.predict(X_test_scaled)
```

```
In [18]: dt_model = DecisionTreeClassifier(random_state=42)
dt_model.fit(X_train_scaled, y_train)
y_pred_dt = dt_model.predict(X_test_scaled)
```

```
In [20]: print("📊 Logistic Regression:\n", classification_report(y_test, y_pred_lr))
print("📊 k-NN:\n", classification_report(y_test, y_pred_knn))
print("📊 Decision Tree:\n", classification_report(y_test, y_pred_dt))
```

```
📊 Logistic Regression:
              precision    recall  f1-score   support

     0       1.00        0.98        0.99         62
     1       0.99        1.00        1.00        138

 accuracy          0.99         200
 macro avg         1.00         0.99         0.99         200
 weighted avg         1.00         0.99         0.99         200
```

```
📊 k-NN:
              precision    recall  f1-score   support

     0       0.95        0.84        0.89         62
     1       0.93        0.98        0.95        138

 accuracy          0.94         200
 macro avg         0.94         0.91         0.92         200
 weighted avg         0.94         0.94         0.93         200
```

```
📊 Decision Tree:
              precision    recall  f1-score   support

     0       0.94        0.95        0.94         62
     1       0.98        0.97        0.97        138

 accuracy          0.96         200
 macro avg         0.96         0.96         0.96         200
 weighted avg         0.97         0.96         0.97         200
```

